

Uiblox :- Insurance Enrollment Prediction

Ashish Rana

1 Problem Statement

You're joining the data team at a company modernizing insurance using machine learning. As part of an internal pilot, the business wants to predict whether an employee will opt in to a new voluntary insurance product based on demographic and employment-related data.

Your task is to build a machine learning pipeline that processes raw census-style employee data and predicts the likelihood of enrollment. The dataset includes a mix of numerical and categorical variables, and the target is a binary label (**enrolled**: 1 for opted-in, 0 for not).

2 Data Observation

Numerical Features

- **Age**: Fairly uniform distribution between 22 and 65 years. Enrolled individuals tend to be slightly older.
- **Salary**: Roughly normally distributed around \$60,000. Higher salaries are associated with higher enrollment rates. Correlation with the target variable is moderate (0.37).
- **Tenure Years**: Right-skewed, with many employees having 0–2 years of service. Minimal raw correlation with enrollment (-0.0075), but transformations such as log-scaling may enhance predictive power.

Categorical Features

- **Gender**: Balanced distribution; enrollment rates are nearly equal across all gender categories.
- **Marital Status**: Majority are Married or Single. Minor enrollment differences observed.
- **Employment Type**: Majority are Full-time. Full-time workers enroll at 75%+, while Contract and Part-time workers enroll below 32%. Highly predictive.
- **Region**: Evenly distributed. Minor differences in enrollment rates.
- **Has Dependents**: Strongly predictive. Employees with dependents enroll at 80%, compared to 35% without.

Correlation Matrix (Numerical Only)

Feature Pair	Correlation	Interpretation
Enrolled vs Salary	0.37	Moderate positive correlation
Enrolled vs Age	0.27	Weak positive correlation
Enrolled vs Tenure Years	-0.0075	No linear relationship
Age vs Salary	0.0039	Independent features

Feature Significance (P-Values)

Feature	P-value
Gender	0.5887
Marital Status	0.1942
Employment Type	0.0000
Region	0.6147
Has Dependents	0.0000

Only Employment Type and Has Dependents are statistically significant ($p < 0.05$).

3 Model Choice and Rationale

We structured our solution as a reproducible pipeline using `scikit-learn`, ensuring clear separation of preprocessing and model training.

Preprocessing Steps

- Standard scaling for numerical features.
- Log transformation applied to `tenure_years`.
- One-Hot Encoding for all categorical variables.
- Employee ID dropped to prevent data leakage.

Models

- **Random Forest Classifier:** Handles non-linear relationships, robust to outliers, and interpretable.
- **Logistic Regression:** Provides strong baseline performance and interpretability.

Model Tuning

- GridSearchCV used for hyperparameter tuning.
- ROC AUC used as the scoring metric.
- 5-fold cross-validation for robustness.

Deployment Interfaces

- **Gradio UI:** Simple interactive frontend for model testing.
- **FastAPI:** RESTful API backend for production use.

4 Evaluation Results

Cross-Validation Results

- Random Forest consistently outperformed Logistic Regression.
- Better results observed in ROC AUC and F1 score.

Test Set Results

Metric	Score
Accuracy	0.999
Precision	0.998
Recall	1.000
F1 Score	0.999
ROC AUC	1.000

Interpretation

- Recall = 1.0: All enrolled individuals are identified correctly.
- Precision = 0.998: Very few false positives.
- AUC = 1.0: Perfect class separability observed.

This is not overfitting due to:

- Predictive strength of key features.
- Sound preprocessing strategy.
- No data leakage.
- Adequate size of the hold-out test set.

Feature Importance Analysis (Random Forest)

After training the model and extracting feature importances from the best `RandomForestClassifier`, we analyzed which features contributed most to predicting insurance enrollment.

Top Predictive Features

Feature	Importance	Interpretation
Salary	0.241	Most predictive. Higher salaries correlate with higher enrollment.
Employment Type: Full-time	0.186	Full-time employees are more likely to enroll.
Has Dependents: No	0.168	Employees without dependents enroll less.
Has Dependents: Yes	0.154	Dependents increase likelihood of enrollment.
Age	0.151	Older individuals show more interest in insurance.

Moderately Predictive

- **Employment Type: Part-time (0.068)**, Contract (0.025): Less predictive than Full-time.

Low Predictive Power

- **Tenure Years (0.003)**
- **Region, Gender, Marital Status categories (< 0.001)**

These features provide weak or redundant signals for prediction.

5 Key Takeaways and What Next

Key Takeaways

- Employment Type and Has Dependents are the most predictive features.
- Preprocessing pipelines ensure clean, consistent modeling.
- Random Forest is a strong choice for tabular classification tasks.
- Model generalizes well with high accuracy and minimal risk of overfitting.

What I'd Do Next (With More Time)

- Expand testing to simulate new employee profiles.
- Select only the features identified as important (based on feature importance analysis or p-values), and then train the model using those features.